

1. Method

We describe the model's architecture in this section, covering the encoder, the decoder, and the training procedure used to fit both components jointly on the benchmark dataset described earlier.

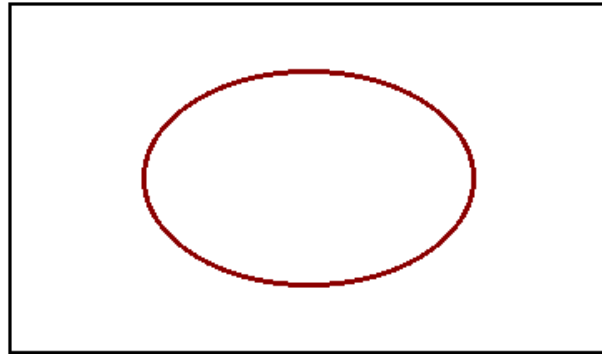


Figure 1: Model overview – the encoder's output feeds the decoder.

Results are discussed further below, along with an ablation study that isolates the contribution of each architectural component.